



加入语雀，获得更好的阅读体验

注册 或 登录 后可以收藏本文随时阅读，还可以关注作者获得最新文章推送

立即加入

## 安装部署（二）

做到这里发现坑了 阿里云的机器无法使用nat表  
不过！用vxlan模型可以搞定（普天同庆）

### flannel

#### flannel的三种网络模型

- host-gw模型：所有node ip必须在同一个物理网关设备下才能使用 它的原理就是：给宿主机添加一个静态路由，指明到达pod之前要经过的宿主机
- Vxlan模型：
- 直接路由模型：当node不在同一个物理网关下，走vxlan模型,在同一个网关下，走host-gw模型

### host-gw模型

下载二进制包

```
1 [root@alice002 ~]# cd /opt/src/
2 [root@alice002 src]# wget "https://github.com/coreos/flannel/releases/download/v0.11.0/flannel-v0.11.0-linux-amd64.tar.gz"
3 [root@alice002 opt]# mkdir flannel-v0.11.0
4 [root@alice002 src]# tar xf flannel-v0.11.0-linux-amd64.tar.gz -C /opt/flannel-v0.11.0
5 [root@alice002 src]# cd /opt/
6 [root@alice002 opt]# ln -s /opt/flannel-v0.11.0 flannel
```

Plain Text

复制代码

拷贝证书到flannel目录

```
1 [root@alice002 opt]# cd flannel
2 [root@alice002 flannel]# mkdir cert
3 [root@alice002 flannel]# cd cert/
4 [root@alice002 cert]# scp alice001:/opt/certs/ca.pem .
5 [root@alice002 cert]# scp alice001:/opt/certs/client.pem .
6 [root@alice002 cert]# scp alice001:/opt/certs/client-key.pem .
7 [root@alice002 cert]# ll
8 total 12
9 -rw-r--r-- 1 root root 1346 Feb  8 20:49 ca.pem
10 -rw----- 1 root root 1679 Feb  8 20:50 client-key.pem
11 -rw-r--r-- 1 root root 1363 Feb  8 20:50 client.pem
```

Plain Text

复制代码

配置文件&启动脚本

Plain Text

复制代码

```
1 [root@alice002 cert]# cd ..
2 [root@alice002 flannel]# vim subnet.env
3 [root@alice002 flannel]# cat subnet.env # 注意网段要修改
4 FLANNEL_NETWORK=172.187.0.0/16
5 FLANNEL_SUBNET=172.187.173.1/24
6 FLANNEL_MTU=1500
7 FLANNEL_IPMASQ=false
8 [root@alice002 flannel]# vi flanneld.sh
9 [root@alice002 flannel]# cat flanneld.sh #注意IP和网卡名称要修改
10 #!/bin/sh
11 ./flanneld
12 --public-ip=172.23.187.173 \
13 --etcd-endpoints=https://172.23.187.175:2379,https://172.23.187.173:2379,https://172.23.187.174:2379 \
14 --etcd-keyfile=./cert/client-key.pem
15 --etcd-certfile=./cert/client.pem
16 --etcd-cafile=./cert/ca.pem
17 --iface=eth0 \
18 --subnet-file=./subnet.env
19 --healthz-port=2401
20 [root@alice002 flannel]# chmod u+x flanneld.sh
```

etcd增加host-gw模型

Plain Text

复制代码

```
1 这里是etcd集群所以在任意一台操作即可
2 [root@alice002 flannel]# cd /opt/etcd/
3 [root@alice002 etcd]# ./etcdctl set /coreos.com/network/config '{"Network": "172.187.0.0/16", "Backend": {"Type": "host-gw"}}'
4 {"Network": "172.187.0.0/16", "Backend": {"Type": "host-gw"}}
5 [root@alice002 etcd]# ./etcdctl get /coreos.com/network/config
6 {"Network": "172.187.0.0/16", "Backend": {"Type": "host-gw"}}
```

启动flannel

Plain Text

复制代码

```
1
2 [root@alice002 etcd]# vim /etc/supervisord.d/flannel.ini
3 [root@alice002 etcd]# cat /etc/supervisord.d/flannel.ini
4 [program:flanneld-alice002]
5 command=/opt/flannel/flanneld.sh
6 numprocs=1
7 directory=/opt/flannel
8 autostart=true
9 autorestart=true
10 startsecs=30
11 startretries=3
12 exitcodes=0,2
13 stopsignal=QUIT
14 stopwaitsecs=10
15 user=root
16 redirect_stderr=true
17 stdout_logfile=/data/logs/flanneld/flanneld.stdout.log
18 stdout_logfile_maxbytes=64MB
19 stdout_logfile_backups=4
20 stdout_capture_maxbytes=1MB
21 stdout_events_enabled=false
22 [root@alice002 etcd]# mkdir -p /data/logs/flanneld/
23 [root@alice002 etcd]# supervisorctl update
24 [root@alice002 etcd]# supervisorctl status
25 etcd-server-alice002          RUNNING    pid 13813, uptime 9 days, 5:03:14
26 flanneld-alice002            RUNNING    pid 10259, uptime 0:00:37
27 kube-apiserver-alice002      RUNNING    pid 14745, uptime 9 days, 2:03:16
28 kube-controller-manager-alice002  RUNNING    pid 14731, uptime 9 days, 2:05:00
29 kube-kubelet-alice002        RUNNING    pid 14901, uptime 9 days, 0:51:26
30 kube-proxy-alice002          RUNNING    pid 26020, uptime 8 days, 11:40:05
31 kube-scheduler-alice002      RUNNING    pid 14344, uptime 9 days, 2:35:35
```

查看路由

Plain Text | 复制代码

```
1 [root@alice002 flannel]# route -n
2 Kernel IP routing table
3 Destination      Gateway            Genmask           Flags Metric Ref    Use Iface
4 0.0.0.0           172.23.255.253    0.0.0.0          UG        0      0      0 eth0
5 169.254.0.0       0.0.0.0           255.255.0.0      U        1002    0      0 eth0
6 172.23.0.0        0.0.0.0           255.255.0.0      U         0      0      0 eth0
7 172.187.173.0     0.0.0.0           255.255.255.0    U         0      0      0 docker0
8 172.187.174.0     172.23.187.174    255.255.255.0    UG        0      0      0 eth0
9 # 要取到172.187.174.0/24 网络的包网关地址是172.23.187.174
```

然而当你设置完发现 还是不通  
这是因为阿里云不支持iptables的nat表  
所以需要用到VxLAN模型

VxLAN模型

停止进程

Plain Text | 复制代码

```
1 [root@alice002 flannel]# supervisorctl status
2 etcd-server-alice002                RUNNING pid 13813, uptime 11 days, 4:31:41
3 flanneld-alice002                   RUNNING pid 30447, uptime 23:07:38
4 kube-apiserver-alice002             RUNNING pid 14745, uptime 11 days, 1:31:43
5 kube-controller-manager-alice002    RUNNING pid 14731, uptime 11 days, 1:33:27
6 kube-kubelet-alice002               RUNNING pid 14901, uptime 11 days, 0:19:53
7 kube-proxy-alice002                 RUNNING pid 26020, uptime 10 days, 11:08:32
8 kube-scheduler-alice002             RUNNING pid 14344, uptime 11 days, 2:04:02
9 [root@alice002 flannel]# supervisorctl stop flanneld-alice002
10 [root@alice002 flannel]# ps -ef |grep flannel
11 root      5571 29646  0 20:50 pts/0    00:00:00 grep --color=auto flannel
12 root      30448    1  0 Feb09 ?        00:00:11 ./flanneld --public-ip=172.23.187.173 --etcd-
13 endpoints=https://172.23.187.175:2379,https://172.23.187.173:2379,https://172.23.187.174:2379 --etcd-keyfile=./cert/client-key.pem --etcd-
14 certfile=./cert/client.pem --etcd-cafile=./cert/ca.pem --iface=eth0 --subnet-file=./subnet.env --healthz-port=2401
15 [root@alice002 flannel]# bash
16 [root@alice002 flannel]# exit
17 [root@alice002 flannel]# kill 30448
18 [root@alice002 flannel]# kill 30448
19 -bash: kill: (30448) - No such process
20 [root@alice002 flannel]# ps -ef |grep flannel
21 root      5695 29646  0 20:51 pts/0    00:00:00 grep --color=auto flannel
```

删除路由

Plain Text | 复制代码

```
1 [root@alice002 flannel]# route -n
2 Kernel IP routing table
3 Destination      Gateway            Genmask           Flags Metric Ref    Use Iface
4 0.0.0.0           172.23.255.253    0.0.0.0          UG        0      0      0 eth0
5 169.254.0.0       0.0.0.0           255.255.0.0      U        1002    0      0 eth0
6 172.23.0.0        0.0.0.0           255.255.0.0      U         0      0      0 eth0
7 172.187.173.0     0.0.0.0           255.255.255.0    U         0      0      0 docker0
8 172.187.174.0     172.23.187.174    255.255.255.0    UG        0      0      0 eth0
9 [root@alice002 flannel]# route del -net 172.187.174.0/24 gw 172.23.187.174
10 [root@alice002 flannel]# route -n
11 Kernel IP routing table
12 Destination      Gateway            Genmask           Flags Metric Ref    Use Iface
13 0.0.0.0           172.23.255.253    0.0.0.0          UG        0      0      0 eth0
14 169.254.0.0       0.0.0.0           255.255.0.0      U        1002    0      0 eth0
15 172.23.0.0        0.0.0.0           255.255.0.0      U         0      0      0 eth0
16 172.187.173.0     0.0.0.0           255.255.255.0    U         0      0      0 docker0
```

删除设置vxlan模型

Plain Text | 复制代码

```
1 [root@alice001 ~]# etcdctl ls coreos.com/network/config/
2 /coreos.com/network/config
3 [root@alice001 ~]# etcdctl rm /coreos.com/network/config
4 PrevNode.Value: {"Network": "172.187.0.0/16", "Backend": {"Type": "host-gw"}}
5
6 [root@alice001 ~]# etcdctl ls coreos.com/network/subnets/
7 /coreos.com/network/subnets/172.187.174.0-24
8 /coreos.com/network/subnets/172.187.173.0-24
9 [root@alice001 ~]# etcdctl rm /coreos.com/network/subnets/172.187.174.0-24
10 Error: x509: certificate signed by unknown authority
11 [root@alice001 ~]# etcdctl rm /coreos.com/network/subnets/172.187.174.0-24
12 PrevNode.Value: {"PublicIP": "172.23.187.174", "BackendType": "host-gw"}
13 [root@alice001 ~]# etcdctl rm /coreos.com/network/subnets/172.187.173.0-24
14 PrevNode.Value: {"PublicIP": "172.23.187.173", "BackendType": "host-gw"}
15 [root@alice001 ~]# etcdctl ls coreos.com/network/subnets/
16
17
18 [root@alice001 ~]# etcdctl set /coreos.com/network/config '{"Network": "172.187.0.0/16", "Backend": {"Type": "VxLAN"}}'
19 {"Network": "172.187.0.0/16", "Backend": {"Type": "VxLAN"}}
20 [root@alice001 ~]# etcdctl get /coreos.com/network/config
21 {"Network": "172.187.0.0/16", "Backend": {"Type": "VxLAN"}}
```

重新启动flannel

Plain Text | 复制代码

```
1 [root@alice002 flannel]# supervisorctl start flanneld-alice002
2 flanneld-alice002: started
3 [root@alice002 flannel]# route -n
4 Kernel IP routing table
5
6 Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
7 0.0.0.0          172.23.255.253 0.0.0.0         UG      0      0      0 eth0
8 169.254.0.0      0.0.0.0        255.255.0.0     U       1002   0      0 eth0
9 172.23.0.0       0.0.0.0        255.255.0.0     U       0      0      0 eth0
10 172.187.173.0    0.0.0.0        255.255.255.0   U       0      0      0 docker0
11 172.187.174.0    172.187.174.0  255.255.255.0   UG      0      0      0 flannel.1
12
13 [root@alice002 flannel]#
14 [root@alice002 flannel]# kubectl get pod -o wide
15
16 NAME          READY   STATUS    RESTARTS   AGE   IP            NODE          NOMINATED NODE   READINESS GATES
17 nginx-ds-2wgnl 1/1     Running   0          4h51m 172.187.173.2 alice002.host.com <none>          <none>
18 nginx-ds-tph79 1/1     Running   0          4h50m 172.187.174.2 alice003.host.com <none>          <none>
19
20 [root@alice002 flannel]# ping 172.187.174.2
21 PING 172.187.174.2 (172.187.174.2) 56(84) bytes of data.
22 64 bytes from 172.187.174.2: icmp_seq=1 ttl=63 time=0.339 ms
23 64 bytes from 172.187.174.2: icmp_seq=2 ttl=63 time=0.296 ms
24 64 bytes from 172.187.174.2: icmp_seq=3 ttl=63 time=0.308 ms
25 ^C
26 --- 172.187.174.2 ping statistics ---
27 3 packets transmitted, 3 received, 0% packet loss, time 2001ms
28 rtt min/avg/max/mdev = 0.296/0.314/0.339/0.023 ms
29
30 [root@alice002 flannel]#
```

直接路由模型

Plain Text | 复制代码

```
1 etcdctl set /coreos.com/network/config '{"Network": "172.7.0.0/16", "Backend": {"Type": "VxLAN", "Directrouting": true}}'
```

snat优化

Plain Text | 复制代码

```
1 ~]# docker run -it --rm nginx:1.7.9 bash
2 root@ad83cb1da6f6:/# tee /etc/apt/sources.list << EOF
3 > deb http://mirrors.163.com/debian/ jessie main non-free contrib
4 > deb http://mirrors.163.com/debian/ jessie-updates main non-free contrib
5 > EOF
6 deb http://mirrors.163.com/debian/ jessie main non-free contrib
7 deb http://mirrors.163.com/debian/ jessie-updates main non-free contrib
8 root@ad83cb1da6f6:/# apt-get update && apt-get install curl -y
9 root@ad83cb1da6f6:/# curl -k https://www.baidu.com
```

Plain Text | 复制代码

```
1 [root@alice001 ~]# docker ps
2 CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS                NAMES
3 ad83cb1da6f6   nginx:1.7.9 "bash"                 4 minutes ago Up 4 minutes   80/tcp, 443/tcp      affectionate_shtern
4 [root@alice001 ~]# docker commit -p ad83cb1da6f6 harbor.od.com/public/nginx:curl
5 sha256:a3cead594cb0674fc8217f9274066d634e07ed85eebed1e754d17d39a819e9d3
6 [root@alice001 ~]# docker push harbor.od.com/public/nginx:curl
```

Plain Text | 复制代码

```
1 ~]# yum install iptables-services -y
2 [root@alice002 ~]# systemctl start iptables
3 [root@alice002 ~]# systemctl enable iptables
4 Created symlink from /etc/systemd/system/basic.target.wants/iptables.service to /usr/lib/systemd/system/iptables.service.
5 [root@alice002 ~]# iptables-save | grep -i postrouting
6 :POSTROUTING ACCEPT [8:486]
7 :KUBE-POSTROUTING - [0:0]
8 -A POSTROUTING -m comment --comment "kubernetes postrouting rules" -j KUBE-POSTROUTING
9 -A POSTROUTING -s 172.187.173.0/24 ! -o docker0 -j MASQUERADE
10 -A KUBE-POSTROUTING -m comment --comment "kubernetes service traffic requiring SNAT" -m mark --mark 0x4000/0x4000 -j MASQUERADE
11 [root@alice002 ~]# iptables -t nat -D POSTROUTING -s 172.187.173.0/24 ! -o docker0 -j MASQUERADE
12 [root@alice002 ~]# iptables -t nat -I POSTROUTING -s 172.187.173.0/24 ! -d 172.187.0.0/16 ! -o docker0 -j MASQUERADE
13 [root@alice002 ~]# iptables-save | grep -i postrouting
14 :POSTROUTING ACCEPT [2:120]
15 :KUBE-POSTROUTING - [0:0]
16 -A POSTROUTING -s 172.187.173.0/24 ! -d 172.187.0.0/16 ! -o docker0 -j MASQUERADE
17 -A POSTROUTING -m comment --comment "kubernetes postrouting rules" -j KUBE-POSTROUTING
18 -A KUBE-POSTROUTING -m comment --comment "kubernetes service traffic requiring SNAT" -m mark --mark 0x4000/0x4000 -j MASQUERADE
19 [root@alice002 ~]# iptables-save | grep -i reject
20 -A INPUT -j REJECT --reject-with icmp-host-prohibited
21 -A FORWARD -j REJECT --reject-with icmp-host-prohibited
22 [root@alice002 ~]# iptables -t filter -D INPUT -j REJECT --reject-with icmp-host-prohibited
23 [root@alice002 ~]# iptables -t filter -D FORWARD -j REJECT --reject-with icmp-host-prohibited
24 [root@alice002 ~]# iptables-save > /etc/sysconfig/iptables
25 [root@alice002 ~]# service iptables save
26 iptables: Saving firewall rules to /etc/sysconfig/iptables:[ OK ]
27 [root@alice002 ~]#
```

### 注意docker重启后操作

docker服务重启后,会再次增加该规则,要注意在每次重启docker服务后,删除该规则

验证:

修改后会影响到docker原本的iptables链的规则,所以需要重启docker服务

Plain Text | 复制代码

```
1 [root@hdss7-21 ~]# systemctl restart docker
2 [root@hdss7-21 ~]# iptables-save |grep -i postrouting|grep docker0
3 -A POSTROUTING -s 172.7.21.0/24 ! -o docker0 -j MASQUERADE
4 -A POSTROUTING -s 172.7.21.0/24 ! -d 172.7.0.0/16 ! -o docker0 -j MASQUERADE
5 # 可以用iptables-restore重新应用iptables规则,也可以直接再删
6 [root@hdss7-21 ~]# iptables-restore /etc/sysconfig/iptables
7 [root@hdss7-21 ~]# iptables-save |grep -i postrouting|grep docker0
8 -A POSTROUTING -s 172.7.21.0/24 ! -d 172.7.0.0/16 ! -o docker0 -j MASQUERADE
```

Plain Text | 复制代码

```
1 [root@alice002 ~]# kubectl get pod -o wide
2 NAME          READY   STATUS    RESTARTS   AGE   IP            NODE          NOMINATED NODE   READINESS GATES
3 nginx-ds-v96dl 1/1     Running   0           21h   172.187.174.2 alice003.host.com <none>          <none>
4 nginx-ds-wn25q 1/1     Running   0           21h   172.187.173.2 alice002.host.com <none>          <none>
5 [root@alice002 ~]# kubectl exec -it nginx-ds-wn25q bash
6 root@nginx-ds-wn25q:/# curl 172.187.174.2
7 <!DOCTYPE html>
8 <html>
9 <head>
10 <title>Welcome to nginx!</title>
11 <style>
12     body {
13         width: 35em;
14         margin: 0 auto;
15         font-family: Tahoma, Verdana, Arial, sans-serif;
16     }
17 </style>
18 </head>
19 <body>
20 <h1>Welcome to nginx!</h1>
21 <p>If you see this page, the nginx web server is successfully installed and
22 working. Further configuration is required.</p>
23
24 <p>For online documentation and support please refer to
25 <a href="http://nginx.org/">nginx.org</a>.<br/>
26 Commercial support is available at
27 <a href="http://nginx.com/">nginx.com</a>.</p>
28
29 <p><em>Thank you for using nginx.</em></p>
30 </body>
31 </html>
32 root@nginx-ds-wn25q:/#
33 [root@alice002 ~]# curl 172.187.174.2 宿主机也curl一下
34
35
36 [root@alice003 ~]# kubectl logs -f nginx-ds-v96dl
37 172.187.173.0 - - [19/Feb/2021:12:37:57 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.38.0" "-"
38
39
40
41 ---
42 172.187.173.2 - - [19/Feb/2021:13:02:26 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.38.0" "-"
43 172.187.173.0 - - [19/Feb/2021:13:03:36 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.29.0" "-"
```

coredns

创建在线yaml

这里我为了可以公网访问 顺便把真实域名解析了一下  
以下内容中的grep.pro都是可以真实访问的  
后续yaml以及dockerfile地址<http://k8s-yaml.grep.pro> <<http://k8s-yaml.grep.pro/>>

```

1 [root@alice001 ~]# vim /etc/nginx/conf.d/k8s-yaml.od.com.conf
2 [root@alice001 ~]# cat /etc/nginx/conf.d/k8s-yaml.od.com.conf
3 server {
4     listen      80;
5     server_name k8s-yaml.od.com k8s-yaml.grep.pro;
6
7     location / {
8         autoindex on;
9         default_type text/plain;
10        root /data/k8s-yaml;
11    }
12 }
13
14 [root@alice001 ~]# vim /var/named/od.com.zone
15 [root@alice001 ~]# cat /var/named/od.com.zone
16 $ORIGIN od.com.
17 $TTL 600      ; 10 minutes
18 @      IN SOA  dns.od.com. dnsadmin.od.com. (
19     2021012906 ; serial
20     10800      ; refresh (3 hours)
21     900        ; retry (15 minutes)
22     604800     ; expire (1 week)
23     86400      ; minimum (1 day)
24 )
25 NS      dns.od.com.
26 $TTL 60 ; 1 minute
27 dns      A      47.243.20.250
28 harbor   A      47.243.20.250
29 k8s-yaml A      47.243.20.250
30 [root@alice001 ~]# systemctl restart named
31 [root@alice001 ~]# dig -t A k8s-yaml.od.com @172.23.187.175 +short
32 47.243.20.250
33 [root@alice001 ~]# nginx -t
34 nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
35 nginx: configuration file /etc/nginx/nginx.conf test is successful
36 [root@alice001 ~]# systemctl restart nginx
37 [root@alice001 ~]# docker pull coredns/coredns:1.6.1
38 1.6.1: Pulling from coredns/coredns
39 c6568d217a00: Pull complete
40 d7ef34146932: Pull complete
41 Digest: sha256:9ae3b6fcac4ee821362277de6bd8fd2236fa7d3e19af2ef0406d80b595620a7a
42 Status: Downloaded newer image for coredns/coredns:1.6.1
43 [root@alice001 ~]# docker images |grep coredns
44 coredns/coredns      1.6.1              c0f6e815079e      18 months ago      42.2MB
45 [root@alice001 ~]# docker tag c0f6e815079e harbor.od.com/public/coredns:v1.6.1
46 [root@alice001 ~]# docker push !$
47 docker push harbor.od.com/public/coredns:v1.6.1
48 The push refers to repository [harbor.od.com/public/coredns]
49 da1ec456edc8: Pushed
50 225df95e717c: Pushed
51 v1.6.1: digest: sha256:c7bf0ce4123212c87db74050d4cbab77d8f7e0b49c041e894a35ef15827cf938 size: 739
52 [root@alice001 ~]# cd /data/
53 [root@alice001 data]# ls
54 docker etcd harbor logs
55 [root@alice001 data]# mkdir /data/k8s-yaml
56 [root@alice001 data]# cd k8s-yaml/
57 [root@alice001 k8s-yaml]# cat /etc/nginx/conf.d/k8s-yaml.od.com.conf
58 server {
59     listen      80;
60     server_name k8s-yaml.od.com;
61
62     location / {
63         autoindex on;
64         default_type text/plain;
65         root /data/k8s-yaml;
66     }
67 }
68 [root@alice001 k8s-yaml]#
69 [root@alice001 k8s-yaml]# mkdir coredns

```

cm.yaml中 forward . 172.23.187.175为 集群的dns地址也是coredns的上游dns地址

svc.yaml中 clusterIP: 192.168.0.2地址为集群中kubelet定义好的集群dns地址

其余的没动 还没弄懂 有同学懂了可以在这里评论一下

Plain Text | 复制代码

```
1 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/coredns/rbac.yaml
2 serviceaccount/coredns created
3 clusterrole.rbac.authorization.k8s.io/system:coredns created
4 clusterrolebinding.rbac.authorization.k8s.io/system:coredns created
5 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/coredns/cm.yaml
6 configmap/coredns created
7 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/coredns/dp.yaml
8 deployment.apps/coredns created
9 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/coredns/svc.yaml
10 service/coredns created
```

验证

Plain Text | 复制代码

```
1 [root@alice002 ~]# dig -t A www.baidu.com @172.23.187.175 +short
2 www.a.shifen.com.
3 www.wshifen.com.
4 104.193.88.123
5 104.193.88.77
6 [root@alice002 ~]# dig -t A alice001.host.com @172.23.187.175 +short
7 47.243.20.250
8 [root@alice002 ~]# dig -t A alice001.host.com @192.168.0.2 +short
9 47.243.20.250
10 [root@alice002 ~]# dig -t A www.baidu.com @192.168.0.2 +short
11 www.a.shifen.com.
12 www.wshifen.com.
13 104.193.88.123
14 104.193.88.77
15
16
17 [root@alice002 ~]# kubectl create deployment nginx-web --image=harbor.od.com/public/nginx:curl_ps
18 deployment.apps/nginx-web created
19 [root@alice002 ~]# kubectl expose deployment nginx-web --port=80 --target-port=80
20 service/nginx-web exposed
21 [root@alice002 ~]# kubectl get svc
22 NAME          TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)    AGE
23 kubernetes    ClusterIP     192.168.0.1     <none>       443/TCP    20d
24 nginx-web     ClusterIP     192.168.208.221 <none>       80/TCP     4s
25 [root@alice002 ~]# kubectl get svc -n kube-system
26 NAME          TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE
27 coredns       ClusterIP     192.168.0.2     <none>       53/UDP,53/TCP,9153/TCP 14m
28 [root@alice002 ~]# dig -t A nginx-web.default.svc.cluster.local @192.168.0.2 +short
29 192.168.208.221
```

traefik



```
1 [root@alice001 coredns]# cd /data/k8s-yaml/
2 [root@alice001 k8s-yaml]# mkdir traefik
3 [root@alice001 k8s-yaml]# cd traefik/
4 [root@alice001 traefik]# docker pull traefik:v1.7.2-alpine
5 v1.7.2-alpine: Pulling from library/traefik
6 4fe2ade4980c: Pull complete
7 8d9593d002f4: Pull complete
8 5d09ab10efbd: Pull complete
9 37b796c58adc: Pull complete
10 Digest: sha256:cf30141936f73599e1a46355592d08c88d74bd291f05104fe11a8bcce447c044
11 Status: Downloaded newer image for traefik:v1.7.2-alpine
12 [root@alice001 traefik]# docker tag traefik:v1.7.2-alpine harbor.od.com/public/traefik:v1.7.2
13 [root@alice001 traefik]# docker push !$
14 docker push harbor.od.com/public/traefik:v1.7.2
15 The push refers to repository [harbor.od.com/public/traefik]
16 a02beb48577f: Pushed
17 ca22117205f4: Pushed
18 3563c211d861: Pushed
19 df64d3292fd6: Pushed
20 v1.7.2: digest: sha256:6115155b261707b642341b065cd3fac2b546559ba035d0262650b3b3bbdd10ea size: 1157
21 [root@alice001 traefik]# ll
22 -rw-r--r-- 1 root root 1100 Jan  9 16:36 ds.yaml
23 -rw-r--r-- 1 root root 330 Jan  9 16:36 ingress.yaml
24 -rw-r--r-- 1 root root 800 Jan  9 16:36 rbac.yaml
25 -rw-r--r-- 1 root root 269 Jan  9 16:36 svc.yaml
26 [root@alice001 traefik]# vi /etc/nginx/conf.d/od.com.conf
27 [root@alice001 traefik]# cat /etc/nginx/conf.d/od.com.conf
28 upstream default_backend_traefik {
29     server 172.23.187.173:81    max_fails=3 fail_timeout=10s;
30     server 172.23.187.174:81    max_fails=3 fail_timeout=10s;
31 }
32 server {
33     server_name *.od.com *.grep.pro;
34     location / {
35         proxy_pass http://default_backend_traefik;
36         proxy_set_header Host      $http_host;
37         proxy_set_header x-forwarded-for $proxy_add_x_forwarded_for;
38     }
39 }
40
41 [root@alice001 traefik]# nginx -t
42 nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
43 nginx: configuration file /etc/nginx/nginx.conf test is successful
44 [root@alice001 traefik]# systemctl reload nginx
45 [root@alice001 traefik]# vim /var/named/od.com.zone
46 [root@alice001 traefik]# cat /var/named/od.com.zone
47 $ORIGIN od.com.
48 $TTL 600      ; 10 minutes
49 @      IN SOA  dns.od.com. dnsadmin.od.com. (
50         2021012907 ; serial
51         10800      ; refresh (3 hours)
52         900        ; retry (15 minutes)
53         604800     ; expire (1 week)
54         86400      ; minimum (1 day)
55     )
56     NS      dns.od.com.
57 $TTL 60 ; 1 minute
58 dns                A      47.243.20.250
59 harbor             A      47.243.20.250
60 k8s-yaml           A      47.243.20.250
61 traefik            A      47.243.20.250
62 [root@alice001 traefik]# systemctl restart named
63 [root@alice001 traefik]# dig -t A traefik.od.com @172.23.187.175 +short
64 47.243.20.250
```

在任意节点上

Plain Text | 复制代码

```
1 [root@alice002 ~]# #---traefik
2 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/traefik/rbac.yaml
3 serviceaccount/traefik-ingress-controller created
4 clusterrole.rbac.authorization.k8s.io/traefik-ingress-controller created
5 clusterrolebinding.rbac.authorization.k8s.io/traefik-ingress-controller created
6 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/traefik/ds.yaml
7 daemonset.extensions/traefik-ingress created
8 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/traefik/svc.yaml
9 service/traefik-ingress-service created
10 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.od.com/traefik/ingress.yaml
11 ingress.extensions/traefik-web-ui created
12 [root@alice002 ~]#
13 [root@alice002 ~]# kubectl get pod -n kube-public
14 No resources found.
15 [root@alice002 ~]# kubectl get pod -n kube-system
16 NAME                                READY   STATUS    RESTARTS   AGE
17 coredns-6b6c4f9648-g6btt          1/1     Running   0           12h
18 traefik-ingress-62gb2              0/1     ContainerCreating   0           35s
19 traefik-ingress-wz2r7              0/1     ContainerCreating   0           35s
20 重启两台机器的docker
21 [root@alice002 ~]# systemctl restart docker.service
22 [root@alice003 ~]# systemctl restart docker.service
23 [root@alice002 ~]# kubectl get pod -n kube-system -o wide
24 NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE                                NOMINATED NODE   READINESS GATES
25 coredns-6b6c4f9648-g6btt          1/1     Running   0           12h   172.187.174.3   alice003.host.com                 <none>            <none>
26 traefik-ingress-62gb2              1/1     Running   0           2m1s  172.187.173.4   alice002.host.com                 <none>            <none>
27 traefik-ingress-wz2r7              1/1     Running   0           2m1s  172.187.174.4   alice003.host.com                 <none>            <none>
```

# dashboard

```
1 [root@alice001 traefik]# docker pull k8scn/kubernetes-dashboard-amd64:v1.8.3
2 v1.8.3: Pulling from k8scn/kubernetes-dashboard-amd64
3 a4026007c47e: Pull complete
4 Digest: sha256:ebc993303f8a42c301592639770bd1944d80c88be8036e2d4d0aa116148264ff
5 Status: Downloaded newer image for k8scn/kubernetes-dashboard-amd64:v1.8.3
6 [root@alice001 traefik]# docker tag k8scn/kubernetes-dashboard-amd64:v1.8.3 harbor.od.com/public/dashboard:v1.8.3
7 [root@alice001 traefik]# docker push !$
8 docker push harbor.od.com/public/dashboard:v1.8.3
9 The push refers to repository [harbor.od.com/public/dashboard]
10 23ddb8cbb75a: Pushed
11 v1.8.3: digest: sha256:ebc993303f8a42c301592639770bd1944d80c88be8036e2d4d0aa116148264ff size: 529
12 [root@alice001 traefik]# mkdir -p /data/k8s-yaml/dashboard && cd /data/k8s-yaml/dashboard
13 [root@alice001 dashboard]# ll
14 total 16
15 -rw-r--r-- 1 root root 1427 Feb 20 19:11 deployment.yaml
16 -rw-r--r-- 1 root root 347 Feb 20 16:23 ingress.yaml
17 -rw-r--r-- 1 root root 610 Feb 20 18:23 rbac.yaml
18 -rw-r--r-- 1 root root 322 Feb 20 16:22 svc.yaml
19 [root@alice001 dashboard]# vim /var/named/od.com.zone
20 [root@alice001 dashboard]# cat /var/named/od.com.zone
21 $ORIGIN od.com.
22 $TTL 600 ; 10 minutes
23 @ IN SOA dns.od.com. dnsadmin.od.com. (
24     2021012909 ; serial
25     10800 ; refresh (3 hours)
26     900 ; retry (15 minutes)
27     604800 ; expire (1 week)
28     86400 ; minimum (1 day)
29 )
30 NS dns.od.com.
31 $TTL 60 ; 1 minute
32 dns A 47.243.20.250
33 harbor A 172.23.187.175
34 k8s-yaml A 47.243.20.250
35 traefik A 47.243.20.250
36 dashboard A 47.243.20.250
37 [root@alice001 dashboard]# systemctl restart named
38 [root@alice001 dashboard]# dig -t A dashboard.od.com @172.23.187.175 +short
39 47.243.20.250
40 [root@alice001 dashboard]# cd /opt/certs/
41 [root@alice001 certs]# openssl req -new -key dashboard.od.com.key -out dashboard.od.com.csr -subj
42 "/CN=dashboard.od.com/C=CN/ST=BJ/L=Beijing/O=Oldb
43 days 3650
44 Signature ok
45 subject=/CN=dashboard.od.com/C=CN/ST=BJ/L=Beijing/O=OldboyEdu/OU=ops
46 Getting CA Private Key
47 [root@alice001 certs]# ll dashboard.od.com.*
48 -rw-r--r-- 1 root root 1196 Feb 20 18:53 dashboard.od.com.crt
49 -rw-r--r-- 1 root root 1005 Feb 20 18:53 dashboard.od.com.csr
50 -rw----- 1 root root 1679 Feb 20 18:53 dashboard.od.com.key
51 [root@alice001 certs]# cd /etc/nginx/
52 [root@alice001 nginx]# mkdir certs
53 [root@alice001 nginx]# cd certs/
54 [root@alice001 certs]# ls
55 [root@alice001 certs]# cp /opt/certs/dashboard.od.com.key .
56 [root@alice001 certs]# cp /opt/certs/dashboard.od.com.crt .
57 [root@alice001 certs]# ll
58 total 8
59 -rw-r--r-- 1 root root 1196 Feb 20 18:57 dashboard.od.com.crt
60 -rw----- 1 root root 1679 Feb 20 18:57 dashboard.od.com.key
61 [root@alice001 certs]# vim /etc/nginx/conf.d/dashborad.conf
62 [root@alice001 dashboard]# cat /etc/nginx/conf.d/dashborad.conf
63 server {
64     listen 80;
65     server_name dashboard.od.com dashboard.grep.pro;
66     rewrite ^(.*)$ https://${server_name}$1 permanent;
67 }
68
69 server {
70     listen 443 ssl;
71     server_name dashboard.od.com dashboard.grep.pro;
72
73     ssl_certificate "certs/dashboard.od.com.crt";
74     ssl_certificate_key "certs/dashboard.od.com.key";
75     ssl_session_cache shared:SSL:1m;
76     ssl_session_timeout 10m;
77     ssl_ciphers HIGH:!aNULL:!MD5;
78     ssl_prefer_server_ciphers on;
79
80     location / {
```

```
80     proxy_pass http://default_backend_traefik;
81     proxy_set_header Host      $http_host;
82     proxy_set_header x-forwarded-for $proxy_add_x_forwarded_for;
83 }
84 }
85 [root@alice001 dashboard]#
86 [root@alice001 dashboard]# nginx -t
87 nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
88 nginx: configuration file /etc/nginx/nginx.conf test is successful
89 [root@alice001 dashboard]# systemctl restart nginx
```

这里部署完后我的docker pull不到镜像了 报错

Plain Text |  复制代码

```
1 [root@alice002 ~]# docker pull narbor.od.com/public/dashboard:v1.10.1
2 Error response from daemon: error parsing HTTP 404 response body: invalid character 'p' after top-level value: "404 page not found\n"
```

把nginx的dashboard.conf删掉后就可以恢复 没有搞明白是什么问题 如果有知道的朋友麻烦评论告诉我一下

节点上

Plain Text |  复制代码

```
1 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.grep.pro/dashboard/rbac.yaml
2 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.grep.pro/dashboard/deployment.yaml
3 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.grep.pro/dashboard/svc.yaml
4 [root@alice002 ~]# kubectl apply -f http://k8s-yaml.grep.pro/dashboard/ingress.yaml
```

登陆dashboard

Plain Text |  复制代码

```

1 [root@alice002 ~]# kubectl get secret -n kube-system
2 error: the server doesn't have a resource type "secret"
3 [root@alice002 ~]# kubectl get secret -n kube-system
4 NAME                                     TYPE                                     DATA   AGE
5 coredns-token-snp8                      kubernetes.io/service-account-token    3       20h
6 default-token-z6pmn                    kubernetes.io/service-account-token    3       21d
7 kubernetes-dashboard-admin-token-pbr2v  kubernetes.io/service-account-token    3       16m
8 kubernetes-dashboard-key-holder         Opaque                                  2       6m19s
9 traefik-ingress-controller-token-t27zn kubernetes.io/service-account-token    3       8h
10 [root@alice002 ~]# kubectl describe secret kubernetes-dashboard-admin-token-pbr2v -n kube-system
11 Name:                                     kubernetes-dashboard-admin-token-pbr2v
12 Namespace:                               kube-system
13 Labels:                                   <none>
14 Annotations: kubernetes.io/service-account.name: kubernetes-dashboard-admin
15              kubernetes.io/service-account.uid: 1f03a210-3dae-4b10-9a19-1c5b6679edd4
16
17 Type: kubernetes.io/service-account-token
18
19 Data
20 ====
21 ca.crt:      1346 bytes
22 namespace:   11 bytes
23 token:
24 eyJhbGciOiJIUzI1NiJ5ZXR1cy5pby9zZXJ2aWNlYWNjb3VudC9uYW1lc3BhY2UiOiJrdWJ1LXN5c3R1bSIsImtYmVybmV0ZXMuaW8vc2Vydm1jZWZlY291bnQvc2VjcmV0Lm5hbWUiOiJrc
25 y23DLvcekzs_wKk7D1KSUDTF_yGF9GnQZ_ECA_4d8yH2q3l0vwpCcItXw0H_Ys0aGw5t8wZbATSUKEEZfjAULXXnZREP9Aa8as14i1tcgw2DGcHxyBCcP9bvH2cj3INsat3lBcmotr3Y3j

```